

Computational Cognitive Modeling

HW 2 Review

03/10/2020

RL Methods

I. Policy Iteration: Evaluation, Improvement

II. Monte Carlo Methods


III. Temporal Difference Learning

i. Dynamic Programming: Policy Iteration

- Two stages
- Policy Evaluation: assign value to states
- Uses your current policy to get values

π


Current (randomized) policy



Bellman equation

$$V^\pi(s) = \sum_a \pi(s, a) \sum_{s'} \mathcal{P}_{ss'}^a [\mathcal{R}_{ss'}^a + \gamma V^\pi(s')]$$

- Policy Improvement: use values to choose actions
 - Uses values of states to choose actions
 - Q-values: state-action pairs



$$\begin{aligned} q_\pi(s, a) &\doteq \mathbb{E}[R_{t+1} + \gamma v_\pi(S_{t+1}) \mid S_t = s, A_t = a] \\ &= \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_\pi(s')]. \end{aligned}$$

i. Policy Iteration

****look at Problem 1 code****

RL Methods

I. Policy Iteration: Evaluation, Improvement

II. Monte Carlo Methods

III. Temporal Difference Learning

ii. First-visit Monte Carlo

- Each time you *initially* visit a state, record the rewards you receive until the end of an episode
- Sample *many episodes*, then average the rewards received for each state
 - This gives you an estimate of the value of the state or action
- Use this estimate to update your policy (greedy method)

ii. First-visit Monte Carlo

****look at Problem 2 code****

****Especially the pseudo-code explanation****

RL Methods

I. Policy Iteration: Evaluation, Improvement

II. Monte Carlo Methods

III. Temporal Difference Learning

iiii. Temporal Difference

- Incremental updating: Q-learning update rule
 - “Off-policy”: Q-value based on max in the next state

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t) \right].$$

iiii. Temporal Difference

- Incremental updating: Q-learning update rule
 - “Off-policy”: Q-value based on max in the next state

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t) \right].$$

This differs from SARSA (on-policy):

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t) \right].$$



iiii. Temporal Difference

- Incremental updating: Q-learning update rule
 - “Off-policy”: Q-value based on max in the next state

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t) \right].$$

- For HW Problem 7: adopt some code from your monte-carlo solution
 - Previously computed discounted average returns from an episode
 - Here: asked to *incrementally* update the Q-value
 - While balancing explore/exploit

iii. Temporal Difference

****look at Problem 7 code****

Questions?